

Redis

"It does not matter how slowly you go as long as you do not stop."
- Confucius



Outline

1. Introduction

2. How Redis Works?

- Why is Redis So Fast?
- How Redis store data?
- Expired Deletion
- Data Persistence


3. In Practices

- Data Structures
- Practices

1. Introduction

1.1. Definition

- Redis is an open-source, **in-memory** database.
- Redis provides **a variety of data structures** to support different business scenarios.
- Redis also **supports many features**: transactions, persistence, Lua scripts, multiple cluster solutions, publish/subscribe mode, memory elimination mechanism, ...
- Use cases: caching, key-value database, message queue, ...




1.2. Redis vs Memcached

- Similarities:
 - In-memory databases and generally used as cache.
 - Expiration policies
 - High performance
- Differences:
 - Redis (single-thread) runs on single core. Memcached runs multiple cores.
 - Redis is better on **read operations and memory efficient**.
Memcached is better on write operations.
> 16 GB impact to the performance of a single Redis instance.
 - Redis supports **more data structures**.
 - Redis supports **data persistence**. Memcached does not.
 - Redis supports **cluster mode**. Memcached does not.
 - Redis provides **other features**: transaction, pub/sub, lua script, ...




1.3.1. Architecture / Master-Slave

- Replication: Data is written to a master instance. The main instance sends copies of those commands to a replica asynchronously
- Pros:
 - Scalable for reads
- Cons:
 - Data inconsistency
 - **Failover. If master is down, which slave would become the new master?**
- To solve failover:
 - Sentinel
 - Cluster




1.3.2. Architecture / Sentinel

- Sentinel mode can monitor the master and slave servers
- In addition, sentinel node serves a role in service discovery
- When master goes down, sentinels node elect 1 node as leader to decide which slave becomes master and configure other slaves to follow the new master.
- Pros:
 - Solve Failover
- Cons:
 - Data inconsistency
 - Operational Overhead
 - **Not scale for writes. All writes go to master**




1.3.3. Architecture / Cluster

- Distributes the data on different servers
- Pros:
 - Solve Failover
 - Reduce the system's dependence on a single master node
 - Scalable for read and write
- Cons:
 - Data inconsistency
 - Operational Overhead
- Choose one of modes:
 - Master-Slave
 - Sentinel
 - **Cluster (recommended)**



1.3.3. Architecture / Cluster

- Hash Slots handle the mapping relationship between data and nodes
- A cluster has a total of **16384 hash slots**
- Even distribution on nodes in a cluster
- **One hash slots can have multiple keys**




2. How Redis Works?

2.1. Why is Redis So Fast?

2.1.1. Redis Bottleneck

- What are **bottlenecks** of Redis?
 - **Memory**
 - **Network bandwidth**
 - **Not CPU**



2.1.2. Why is Redis So Fast?

- **In-Memory**
- **Single Thread Model**
 - CPU is not the bottleneck
 - Avoid context switching for multi-threads
 - Multi-threading app require locks or other synchronization mechanisms
→ complexity, bug prone → difficult to gain performance
- **Multiplexing I/O**
 - One thread processes multiple IO streams
- **Efficient Data Structures**

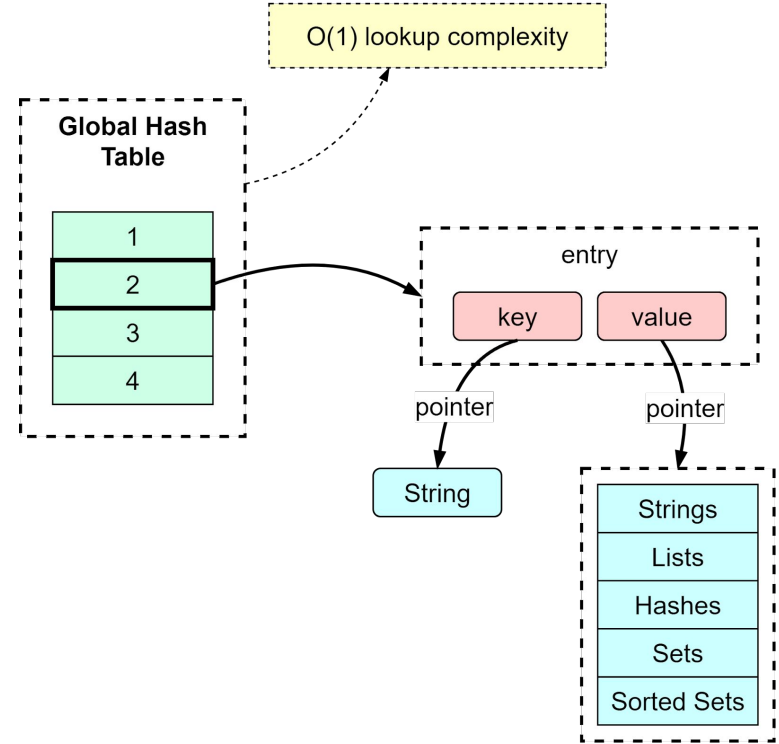
2.1.3. Note

- How to maximize CPU usage?
- Run multiple instances of Redis in the same server/machine
- After version 6.0, Redis is not single-threaded actually
 - Main thread execute commands
 - Other threads handle data persistence, network I/O

2.2. How Redis Store Data?

2.2. How Redis Store Data?

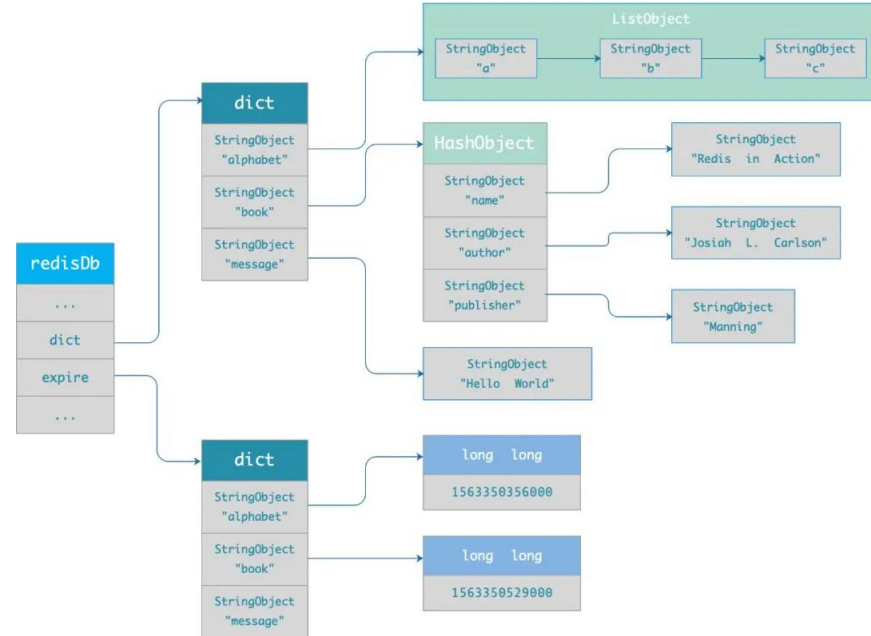
- A dictionary is a hash table → quick search $O(1)$



2.3. Expired Deletion

2.3.1. How to determine if the key has expired?

- 2 dictionaries:
 - Key dictionary
 - Expire dictionary
- When accessing a key, Redis first check if the key exists in the expire dictionary
 - If not, read the key value normally
 - If it exists, then compare with current time.
 - If it is smaller than the current time
 - expired
 - return null




2.3.2. What are the expiration deletion strategies?

- **Passive Way:**

- Key is not deleted immediately right after expiration.
- When the key is accessed, if it expires, then delete the key asynchronously.

→ Problem: some keys is never accessed after they are created

→ memory space wasted



2.3.2. What are the expiration deletion strategies?

- **Active Way:** Periodic job (activeExpireCycle): 10 times/s
 - Randomly select 20 keys from the expired dictionary
 - Check whether these 20 keys have expired and delete the expired keys
 - If the number of expired keys exceeds 25% of the number of randomly selected keys, continue to repeat step 1. If it does not, then the current job stops and wait for the next round

→ Problem: if the active job is long, active job will block other requests.

→ Set timeout for active jobs

2.3.2. What are the expiration deletion strategies?

- **Combine 2 ways: passive way + active way**
- **Problem:** Finding expired keys is not effective
- **New Approach:** Expiration of keys is stored in a Sorted Set (ZSET)
→ find expired keys more effectively

What happens when Redis memory is full?

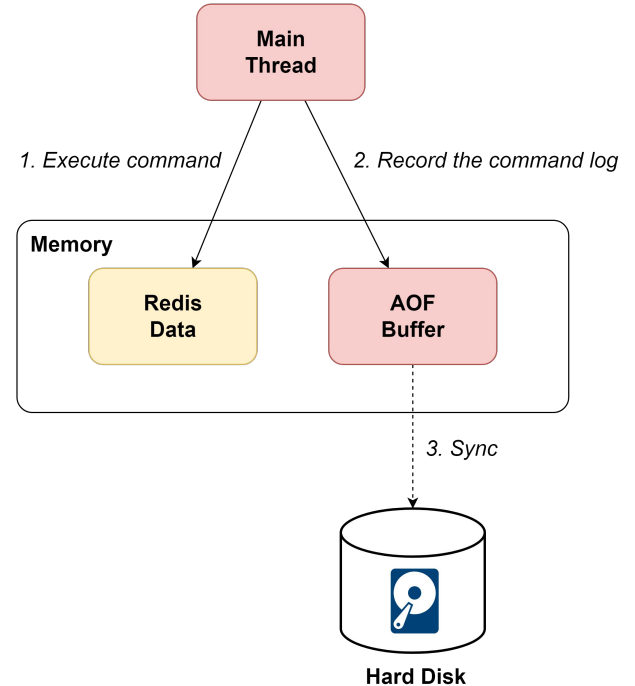
2.4. Memory Eviction

- Noeviction: The default memory eviction
- Random
- TTL: prioritize the elimination of key values that expire earlier
- LRU (Least Recently Used)
- LFU (Least Frequently Used)

2.4. Data Persistence

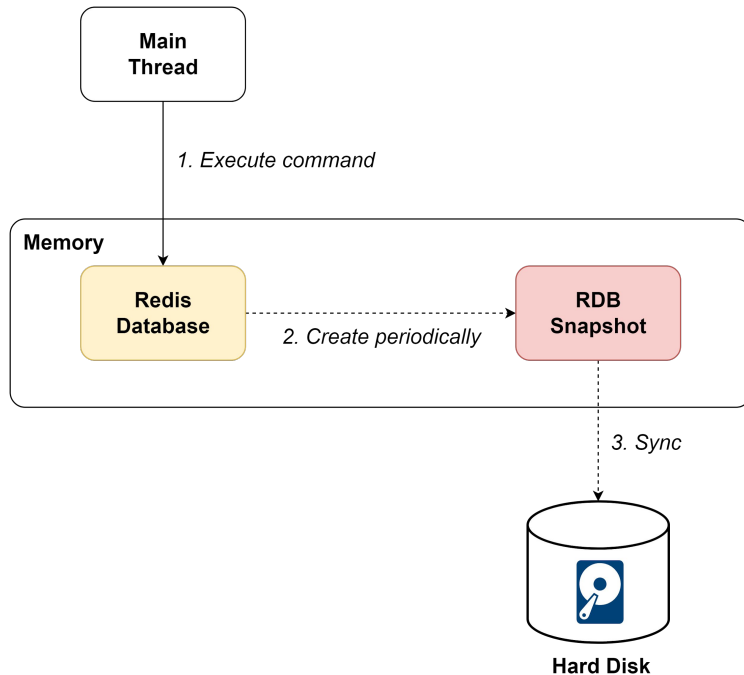
2.4.1. AOF (Append Only File)

- Pros:
 - Avoid additional checking overhead
 - Do not block the execution of the current write operation command
 - Less data loss
- Cons:
 - Data loss
 - May **block other operations**
 - If too many AOF logs → **slow recovery**



2.4.2. RDB (Redis Database)

- **The default mode**
- 2 ways to generate RDB files:
 - Save command: RDB is generated in the main thread → blocking the main thread
 - Bgsave command: a child process generates RDB → avoid blocking the main thread
- Pros:
 - Fast Recovery
- Cons:
 - **Data loss**
 - **A relatively heavy operation**



2.4.3. Hybrid

- Pros:
 - Less data loss
 - Fast recovery
- Cons:
 - Poor readability
 - Poor compatibility. Hybrid persistence AOF file cannot be used in versions prior to Redis 4.0



Hybrid Persistence

3. In Practices

3.1. Data Structures

3.0. Question

- User object: {uid, username, email, age}
- Requirements:
 - C1: access username only
 - C2: update age only
 - C3: get all fields
- Which data structure should we apply for the above 3 cases?
 - Hash
 - String Json

3.1.1. String

- Implementation: SDS (Simple Dynamic String)
 - **SDS can save not only text data, but also binary data**
 - SDS is **safe**, concatenating strings will not cause buffer overflow
- Applications
 - Cache objects (JSON)
 - Count
 - Share session
 - Distributed Lock*

3.1.2. List

- List is a simple list of strings, sorted in insertion order
- Elements can be added to the head or tail of List.
- Applications
 - Store list of elements
 - Message Queue
- Limitation:
 - The maximum length of List is $2^{32} - 1$
 - List does not support multiple consumers
 - Weak order preservation

3.1.3. Hash


- Implementation: **hash table**
- Applications
 - Storing objects if frequently access or update attributes in objects
 - Shopping cart
 - Add item: `HSET cart:{user_id} {product_id} 1`
 - Increase Quantity: `HINCRBY cart:{user_id} {product_id} 1`
 - Total number of items: `HLEN cart:{user_id}`
 - Delete item: `HDEL cart:{user_id} {product_id}`
 - Get an item: `HGET cart:{user_id} {product_id}`
 - Get all items in the shopping cart: `HGETALL cart:{user_id}`

3.1.4. Set

- Set is an unordered and unique set of key values, and its storage order will not be stored in the order of insertion.
- Implementation: **hash table**
- Applications
 - Deduplication of data and ensuring the uniqueness
 - Function on sets: difference, union and intersection
 - Likes, common followers, ...
- Limitations:
 - Max: $2^{32}-1$ elements
 - The calculation complexity of functions on sets is relatively high. → block the main thread

3.1.5. Sorted Set (ZSET)

- Set is an ordered and unique set of key values, and its storage order will not be stored in the order of insertion.
- Implementation: **Skip List**
- Applications
 - Sorting
 - Ranking, Leaderboard



3.2. Practices

3.2.1. How to implement a delay task?

- Context:
 - When ordering a taxi, if there is no car owner to take the order within 10 minutes
 - The platform will cancel your order automatically
 - Remind you that there is no car owner to take the order at the moment
- **Use an sorted set (ZSet)**
- Add element:
 - Value: order_id
 - Score: expiration time
- Poll to get expires order


Key	Value	Score (timestamp)
o1	u1	1704790435
o2	u34	1704793232
o3	u88	1704797344

3.2.2. How to deal with big keys in Redis?

- Big key:
 - The value of type String is greater than 10 KB
 - The number of elements of Hash, List, Set, and ZSet types exceeds 5000
- Effects:
 - Time-consuming → timeout
 - Network congestion
 - Block the worker thread when deleting
 - Memory is unevenly distributed
- Approach:
 - Use the SCAN, HLEN, SCARD, MEMORY USAGE command to find the big key
 - Do not use DEL
 - Recommended: [UNLINK](#). Asynchronous Deletion: Unlink + configure


3.2.3. How to put different keys on a node?

- Limitation:
 - Functions on 2 sets on 2 different nodes → does not work
 - 2 sets in the same hash slot → works
- Use Hashtag
 - Key 1: User_profile:{34}
 - Key 2: User_session:{34}



3.2.4. How to implement distributed locks?

- Problem: Race condition
 - Context: 2 requests concurrently
 - Expectation: **1 request is processed at a time**
- Use SETNX
- Process:
 - SET lock_key value NX PX 10000
 - DEL (UNLINK) lock_key



3.2.5. Best Practices


- Run Redis cluster
- Leverage multi get, pipelines
- Avoid long-run tasks
- Normally set short TTL
- Distribute TTL → avoid thundering herd
- Pick right data structures
- Leverage hashtag
- Understand the time complexity of each command
- [Performance Tuning Best Practices](#)
- ...

Recap

- A cluster has a total of 16384 hash slots. Data is distributed into hash slots.
- CPU is not the bottleneck. Single thread model makes sense
- Leverage data structures and features. Specially care about distributed lock.

Homework

- Implement Distributed Lock
 - 2 requests booking the same seat (A34_S012C) on a flight (FA634)
 - No need to implement DB and booking logic



References

- <https://developer.redis.com/howtos/antipatterns/>
- <https://redis.com/blog/7-redis-worst-practices/>
- <https://architecturenotes.co/redis/>
- <https://medium.com/software-design/why-software-developers-should-care-about-cpu-caches-8da04355bb8a>

Thank you 🙏



Thank you 🙏

