

# Transaction

*“Stop thinking, and end your problems.”*  
- Lao Tzu



# Outline

## 1. Introduction

- Definition
- ACID
- Isolation Levels


## 2. How Transactions Work?

## 3. Common Problems

# 1. Introduction

# 1.1. Definition

- A transaction is a **group of SQL queries** that are treated **atomically**, as a single unit of work.
- All or Nothing
- Syntax:
  - START TRANSACTION / BEGIN: start transaction
  - a group of queries
  - COMMIT: Apply changes
  - ROLLBACK: Discard the changes



## 1.2. ACID

- ACID are properties of transaction in SQL DB
- **Atomicity**: a group of SQL queries is like a query. **all or nothing**.
- **Consistency**:
  - Data should be valid according to all predefined rules.
  - The database should always move from one consistent state to the another, **maintaining all predefined rules, including SQL constraints, business rule**, data types
- **Isolation**:
  - Multiple concurrent transactions do not affect each others.
  - The results of a transaction are usually **invisible** to other transactions until the transaction is complete.
- **Durability**: Committed data would be not lost, even after power failure.

## 1.3. Isolation Levels

## 1.3.1. Read Uncommitted

- Transactions can view the results of **uncommitted transactions**
- Problem: **Dirty Read** → Buggy
- Rarely used in practice
- Use case: A high-frequency logging system
  - Requirement: **performance is prioritized over the accuracy** of the log data

## 1.3.2. Read Committed

- Transactions can view the results of **committed transactions**
- Solve the Dirty Read problem
- In **Postgresql**, Read Committed is the **default** isolation level
- Use case: financial system
  - Requirement: each transaction must only see committed data to prevent anomalies but can tolerate non-repeatable reads.
- Problem: **Read Skew** (Non-repeatable read) - get different value on re-read of a row if another concurrent transaction updates the same row and commits.



## 1.3.3. Repeatable Read

- Repeatable Read guarantees that any rows a transaction reads will “look the **same**” in **subsequent reads within the same transaction**
- Solve the non-repeatable read problem
- In **MySQL**, Repeatable Read is the **default** isolation level
- Problem: **Phantom Read** - get different rows after re-execution of a range query if another transaction adds or removes some rows in the range and commits.
  - A new row appears without knowing where it comes from. It sounds like a phantom.
- Use case: An online retail system
  - Requirement: maintain a consistent view of item prices and availability throughout the session

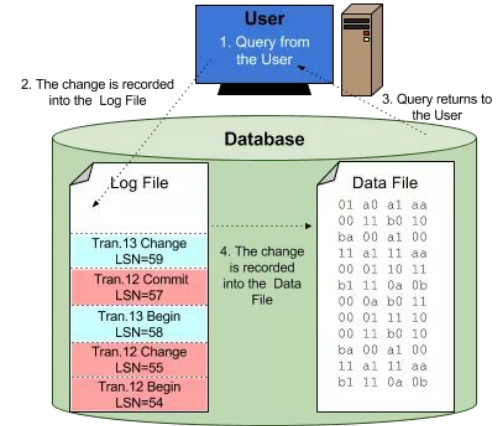
## 1.3.4. Serializable

- SERIALIZABLE solves the phantom read problem by forcing transactions to be ordered so that they can't possibly conflict
- In a nutshell, **SERIALIZABLE places a lock on every row it reads**
- Problem: **Decreasing concurrency**
- Serializable is rarely used in practice.
- Use case: real estate system
  - Requirement: ensures full isolation, preventing phantom reads

## 2. How Transactions Work?

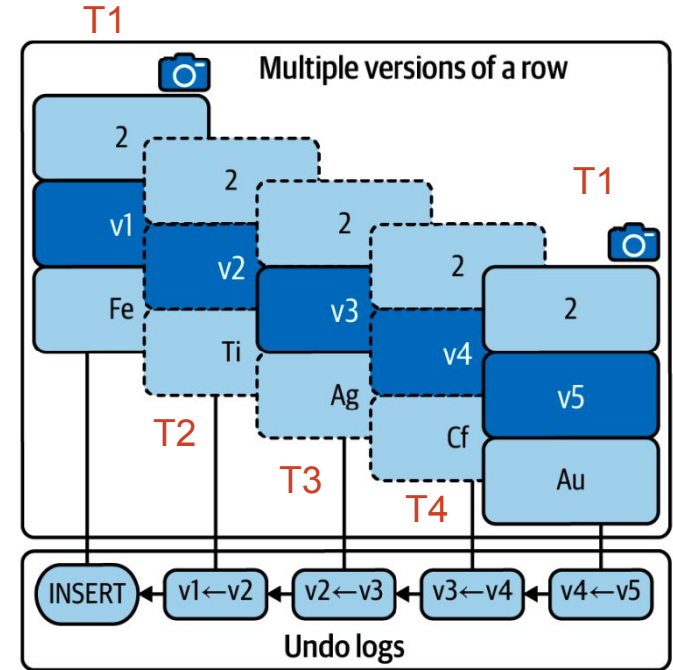
## 2.1. Transaction Logging

- Write Operation:
  - **Write-ahead logging:** The storage engine can then write a record of the change to the transaction log on disk (durable, sequential I/O).
  - The storage engine can change its **in-memory** copy of the data. This is very fast.
  - Later, storage engine update the real data on disk.



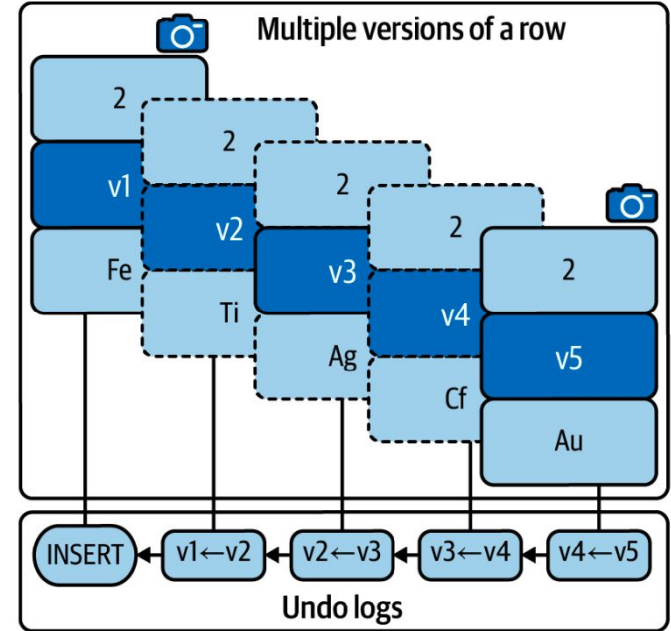
## 2.2. Multiversion Concurrency Control (MVCC)

- InnoDB uses multiversion concurrency control (MVCC) and undo logs to accomplish the A, C, and I properties of ACID
- Multiversion concurrency control means that **changes to a row create a new version** of the row.  
A version of the row is called Undo log.
- **Undo logs record how to reconstruct old row versions and used on REPEATABLE READ and ROLLBACK**
- A snapshot is created when SELECT



## 2.2. Multiversion Concurrency Control (MVCC)

- When the original transaction commits and no other active transaction are holding old snapshot, MySQL purges the related undo logs
- **Undo logs are saved in the InnoDB buffer pool.** Since undo logs reside in buffer pool pages, they use memory and are periodically flushed to disk.
- A transaction can hold many locks and undo logs → **performance impact**



## 2.2. Multiversion Concurrency Control (MVCC)

- [MVVC in Postgres](#)

## 2.3. History list length (HLL)

- History list length (HLL) gauges the amount of old row versions not purged or flushed.
- Monitor and keep HLL short

HLL	Response time (ms)	Baseline increase (%)
0	0.200 ms	
495	0.612 ms	206%
1,089	1.012 ms	406%
2,079	1.841 ms	821%
5,056	3.673 ms	1,737%
11,546	8.527 ms	4,164%



### 3. Common Problems

## 3.1. Large Transaction

- Problem: A transaction modifies too many rows
- Solution:
  - Find the large transaction query
  - **Why it's modifying too many rows?**
  - Change them to modify fewer rows
  - Note: depend business logic

## 3.2. Long-Running Transactions

- Problem: A transaction takes too long to complete
- Causes:
  - The queries in the transaction are too slow
  - App executes too many queries in a transaction
- Solution:
  - (1) Find the root cause. **Maybe data locks**
  - (2) Limit fewer queries in the transaction

## 3.3. Stalled Transaction

- Problem: A transaction is waiting too long between queries in a transaction
- For example:

BEGIN;

SELECT FROM table1 ... ;

– App takes it too long here

SELECT FROM table2 ...;

COMMIT;

- Solution:
  - Depending on app logic
  - **Do these queries need to be a transaction?**
  - Can we use READ COMMITTED to disable gap locking?

## 3.4. Abandoned Transaction

- Problem: Client connection vanished during active transaction
- Causes:
  - App connection leaks
  - Half-closed connections: rarely
- Solution:
  - Fix app logic
  - Generate a report to find abandoned transactions

## 3.5. Notes

- Set transaction isolation level apply the next transaction only. After the next transaction, subsequent transactions use the default transaction isolation level.
- In MySQL, Autocommit is enabled by default.
- Alert on History List Length
- ...

# Recap

- Each isolation level solves the problem of the previous level.  
Read Committed and Repeatable Read are used in practice.
- Transaction works based on MVCC and undo logs. Optimize transaction.
- Transaction and Locking are 2 different mechanism solving 2 different problems

# Homework

- Replay Isolation Level script





# References

- <https://blog.lawrencejones.dev/state-machines/index.html>
- <https://vladmihalcea.com/optimistic-vs-pessimistic-locking/>
- <https://faculty.kutztown.edu/schwesin/spring2022/csc343/lectures/Locking.pdf>

Thank you 🙏

